

- 1. Write a OpenMP program to sort an array on n elements using both sequential and parallel mergesort (using Section). Record the difference in execution time.**

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#define MAX 100000
void merge(int arr[], int l, int m, int r) {
    int i, j, k;
    int n1 = m - l + 1;
    int n2 = r - m;
    int *L = (int *)malloc(n1 * sizeof(int));
    int *R = (int *)malloc(n2 * sizeof(int));
    for (i = 0; i < n1; i++) L[i] = arr[l + i];
    for (j = 0; j < n2; j++) R[j] = arr[m + 1 + j];
    i = 0; j = 0; k = l;
    while (i < n1 && j < n2) {
        arr[k++] = (L[i] <= R[j]) ? L[i++] : R[j++];
    }
    while (i < n1) arr[k++] = L[i++];
    while (j < n2) arr[k++] = R[j++];
    free(L); free(R);
}
void sequentialMergeSort(int arr[], int l, int r) {
    if (l < r) {
        int m = (l + r) / 2;
        sequentialMergeSort(arr, l, m);
        sequentialMergeSort(arr, m + 1, r);
        merge(arr, l, m, r);
    }
}
void parallelMergeSort(int arr[], int l, int r) {
    if (l < r) {
        int m = (l + r) / 2;
        #pragma omp parallel sections
        {
            #pragma omp section
            parallelMergeSort(arr, l, m);
            #pragma omp section
            parallelMergeSort(arr, m + 1, r);
        }
        merge(arr, l, m, r);
    }
}
void copyArray(int *src, int *dest, int n) {
    for (int i = 0; i < n; i++) dest[i] = src[i];
```

```
}

int main() {
    int n;
    int *arr_seq, *arr_par;
    printf("Enter number of elements: ");
    scanf("%d", &n);
    if (n > MAX) {
        printf("Max limit exceeded (%d)\n", MAX);
        return 1;
    }
    arr_seq = (int *)malloc(n * sizeof(int));
    arr_par = (int *)malloc(n * sizeof(int));
    // Generate random data
    for (int i = 0; i < n; i++) {
        arr_seq[i] = rand() % 10000;
    }
    copyArray(arr_seq, arr_par, n);
    double start, end;
    // Sequential sort
    start = omp_get_wtime();
    sequentialMergeSort(arr_seq, 0, n - 1);
    end = omp_get_wtime();
    printf("Sequential MergeSort Time: %.6f seconds\n", end - start);
    // Parallel sort
    start = omp_get_wtime();
    #pragma omp parallel
    {
        #pragma omp single
        parallelMergeSort(arr_par, 0, n - 1);
    }
    end = omp_get_wtime();
    printf("Parallel MergeSort Time: %.6f seconds\n", end - start);
    free(arr_seq); free(arr_par);
    return 0;
}
```

Step 1: Install GCC with OpenMP Support

```
sudo apt update  
sudo apt install gcc
```

Ensure OpenMP is supported:

```
gcc --version
```

Step 2: Save the Program: Save the code above into a file named **mergesort_openmp.c**.

Step 3: Compile with OpenMP: **gcc -fopenmp -o mergesort_openmp mergesort_openmp.c**

Step 4: Run the Program: **./mergesort_openmp**

Output:

Enter number of elements: 100000
Sequential MergeSort Time: 0.092385 seconds
Parallel MergeSort Time: 0.039472 seconds

Viva Questions:

• What is the primary goal of this program?

The program aims to compare the execution time of merge sort implemented sequentially and in parallel using OpenMP to highlight the performance benefits of parallel processing.

• Why is merge sort suitable for parallelization?

Merge sort is a divide-and-conquer algorithm that naturally splits the array into independent subproblems, which can be sorted concurrently before merging, making it highly suitable for parallelization.

• How is parallelism achieved in this program?

Using #pragma omp parallel sections, the two recursive calls of merge sort (left and right sub-arrays) are executed in parallel sections when performing parallelMergeSort.

• What is the purpose of the #pragma omp single directive in the main function?

It ensures that only a single thread initially invokes the recursive parallel merge sort function, preventing redundant recursive calls by all threads.

• How is memory managed during the merge operation?

Temporary arrays (L and R) are dynamically allocated using malloc during each merge and freed after use to prevent memory leaks.

• What would happen if #pragma omp parallel sections were not used in parallelMergeSort?

The program would behave like a sequential merge sort since both recursive calls would execute serially, defeating the purpose of parallelization.

• Why is array copying done before performing parallel sort?

To ensure a fair comparison between sequential and parallel executions by using the same input data, as sorting modifies the array.

• How does OpenMP handle task creation internally for the sections?

Each #pragma omp section is treated as a separate task, potentially executed by different threads from the OpenMP thread pool concurrently.

• What are the limitations of parallel merge sort in this implementation?

Excessive thread creation for small sub-arrays can cause overhead. Also, if the number of recursive levels exceeds available threads, parallelism may not yield significant performance gains.

• How does the use of omp_get_wtime() aid in evaluating performance?

It provides accurate wall-clock time measurement, allowing objective comparison of execution durations for sequential and parallel versions.

2. Write an OpenMP program that divides the Iterations into chunks containing 2 iterations, respectively (OMP_SCHEDULE=static,2). Its input should be the number of iterations, and its output should be which iterations of a parallelized for loop are executed by which thread.

For example, if there are two threads and four iterations, the output might be the following:

- a. Thread 0: Iterations 0 -- 1**
- b. Thread 1: Iterations 2 -- 3**

```
#include <stdio.h>
#include <omp.h>
int main() {
    int num_iterations;
    printf("Enter the number of iterations: ");
    scanf("%d", &num_iterations);

    // Optional: Set number of threads (or use OMP_NUM_THREADS)
    omp_set_num_threads(2);

    printf("\nUsing schedule(static,2):\n\n");
    #pragma omp parallel
    {
        int tid = omp_get_thread_num();
        #pragma omp for schedule(static, 2)
        for (int i = 0; i < num_iterations; i++) {
            printf("Thread %d : Iteration %d\n", tid, i);
        }
    }
    return 0;
}
```

Step 1: Save the Program : Save the code above into a file named **static_schedule.c**

Step 2: Compile with OpenMP: **gcc -fopenmp -o static_schedule static_schedule.c**

Step 3: Run the Program: **./static_schedule**

Output:

Enter the number of iterations: 6

Using schedule(static,2):

```
Thread 0 : Iteration 0
Thread 0 : Iteration 1
Thread 1 : Iteration 2
Thread 1 : Iteration 3
Thread 0 : Iteration 4
Thread 0 : Iteration 5
```

Viva Questions:

- **What does OMP_SCHEDULE=static,2 mean?**
It assigns fixed-size chunks of 2 iterations to each thread in order.
- **What is the purpose of omp parallel for?**
It parallelizes the for loop across multiple threads.
- **How does static scheduling work in OpenMP?**
Iterations are divided before execution and assigned evenly to threads.
- **What is a “chunk”?**
A group of consecutive loop iterations assigned to a thread.
- **What happens if the number of iterations is not a multiple of the chunk size?**
The remaining iterations are still assigned, possibly unevenly.
- **What is omp_get_thread_num() used for?**
It returns the ID of the thread executing the current block.
- **How many threads are created if we use -np 2?**
Two threads will execute the parallel region.
- **Why is output sometimes unordered in parallel loops?**
Because threads execute concurrently and write to output at different times.
- **How does this program demonstrate scheduling clearly?**
It prints which iterations each thread is executing, showing chunk assignments.
- **Can chunk size affect performance?**
Yes, inappropriate chunk sizes can lead to load imbalance or overhead.

3. Write a OpenMP program to calculate n Fibonacci numbers using tasks.

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
// Recursive Fibonacci using OpenMP tasks
int fib(int n) {
    int x, y;
    if (n <= 1)
        return n;
    #pragma omp task shared(x)
    x = fib(n - 1);
    #pragma omp task shared(y)
    y = fib(n - 2);
    #pragma omp taskwait
    return x + y;
}
int main()
{
    int n;
    printf("Enter the number of Fibonacci numbers to calculate: ");
    scanf("%d", &n);
    if (n <= 0) {
        printf("Please enter a positive integer.\n");
        return 0;
    }
    printf("First %d Fibonacci numbers using OpenMP tasks:\n", n);
    double start = omp_get_wtime();
    #pragma omp parallel
    {
        #pragma omp single
        {
            for (int i = 0; i < n; i++) {
                #pragma omp task firstprivate(i)
                {
                    int result = fib(i);
                    #pragma omp critical
                    printf("Fib(%d) = %d\n", i, result);
                }
            }
        }
    }
    double end = omp_get_wtime();
    printf("Execution time: %.6f seconds\n", end - start);
    return 0;
}
```

Step 1: Save the Program : Save the code above into a file named **fib_tasks.c**

Step 2: Compile with OpenMP: **gcc -fopenmp -o fib_tasks fib_tasks.c**

Step 3: Run the Program: **./ fib_tasks.c**

Enter the number of Fibonacci numbers to calculate: 10

First 10 Fibonacci numbers using OpenMP tasks:

```
Fib(0) = 0
Fib(1) = 1
Fib(2) = 1
Fib(3) = 2
Fib(4) = 3
Fib(5) = 5
Fib(6) = 8
Fib(7) = 13
Fib(8) = 21
Fib(9) = 34
Execution time: 0.001234 seconds
```

Viva Questions

- **What is the purpose of using omp task in this Fibonacci program?**
omp task is used to create separate tasks for recursive Fibonacci calls so that they can be executed in parallel, exploiting task-level parallelism.
- **Why is omp taskwait used in the fib function?**
omp taskwait ensures that both recursive tasks (fib(n-1) and fib(n-2)) complete before computing the final result $x + y$.
- **What is the role of the firstprivate(i) clause in the task creation?**
firstprivate(i) ensures each task receives its own private copy of the loop variable i, preserving its value at the time the task was created.
- **Why is omp single used in the main function?**
omp single ensures that only one thread (from the parallel region) creates the set of tasks in the loop, avoiding redundant task creation.
- **What happens if we remove the omp parallel directive in main()?**
Without omp parallel, no threads will be available to execute the created tasks in parallel; the program will execute serially.
- **What is the purpose of omp critical in this program?**
omp critical prevents concurrent writes to the output stream (printf), avoiding mixed or corrupted output from multiple threads.
- **Is this program efficient for large values of n? Why or why not?**
No, it is inefficient for large n due to redundant recursive calls and exponential time complexity.
It demonstrates tasking, not optimized computation.
- **How does task parallelism differ from data parallelism in OpenMP?**
Task parallelism divides the program into independent tasks (e.g., recursive calls), while data parallelism splits data into chunks and applies the same operation.
- **How many tasks are generated when calculating fib(n) using this method?**
Approximately 2^n tasks are generated because each Fibonacci call spawns two subtasks, leading to exponential task creation.
- **What optimization could you apply to improve this Fibonacci task-based program?**
Memoization (caching results) or switching to an iterative Fibonacci approach would significantly reduce redundant computations.

- 4. Write a OpenMP program to find the prime numbers from 1 to n employing parallel for directive. Record both serial and parallel execution times.**

```
#include <stdio.h>
#include <stdlib.h>
#include <omp.h>
#include <math.h>

// Function to check if a number is prime
int is_prime(int num) {
    if (num < 2) return 0;
    if (num == 2) return 1;
    if (num % 2 == 0) return 0;
    for (int i = 3; i <= sqrt(num); i += 2)
        if (num % i == 0)
            return 0;
    return 1;
}

int main() {
    int n;
    printf("Enter the upper limit (n): ");
    scanf("%d", &n);

    if (n < 2) {
        printf("There are no prime numbers less than %d.\n", n);
        return 0;
    }

    printf("\n--- Serial Execution ---\n");
    double start_serial = omp_get_wtime();
    int serial_count = 0;
    for (int i = 2; i <= n; i++) {
        if (is_prime(i)) {
            serial_count++;
            // printf("%d ", i); // Optional: Uncomment to display primes
        }
    }
    double end_serial = omp_get_wtime();
    printf("Total primes found (serial): %d\n", serial_count);
    printf("Serial Execution Time: %.6f seconds\n", end_serial - start_serial);

    printf("\n--- Parallel Execution (OpenMP) ---\n");

    // Set number of threads explicitly
    omp_set_num_threads(4);

    double start_parallel = omp_get_wtime();
    int parallel_count = 0;
```

```
#pragma omp parallel for reduction(+:parallel_count)
for (int i = 2; i <= n; i++) {
    if (is_prime(i)) {
        parallel_count++;
        // #pragma omp critical
        // printf("%d ", i); // Optional: Uncomment to display primes
    }
}
double end_parallel = omp_get_wtime();
printf("Total primes found (parallel): %d\n", parallel_count);
printf("Parallel Execution Time: %.6f seconds\n", end_parallel - start_parallel);

return 0;
}
```

Step 1: Save the Program : prime_parallel.c

Step 2: Compile with OpenMP: gcc -fopenmp -o prime_parallel prime_parallel.c -lm

Step 3: Run the Program: ./ prime_parallel.c

Output:

Enter the upper limit (n): 100000

--- Serial Execution ---

Total primes found (serial): 9592
Serial Execution Time: 0.188745 seconds

--- Parallel Execution (OpenMP) ---

Total primes found (parallel): 9592
Parallel Execution Time: 0.062341 seconds

Viva Questions:

- **What is the primary objective of this OpenMP program?**

The objective is to compare the performance of serial and parallel implementations for counting prime numbers up to a given limit n using OpenMP.

- **Why is sqrt(num) used in the is_prime() function?**

To optimize prime checking by reducing the number of iterations. A number greater than 1 is non-prime if it has a factor less than or equal to its square root.

- **What is the purpose of the directive #pragma omp parallel for reduction (+:parallel_count)?**

It parallelizes the loop and uses a reduction clause to safely accumulate the number of primes (parallel_count) across threads without race conditions.

- **Why is the critical section commented out in the parallel version?**

It was optionally used to safely print prime numbers during parallel execution. However, printing inside parallel regions can slow down performance, so it's often disabled.

- **What are the advantages of using `omp_get_wtime()` in this program?**

`omp_get_wtime()` provides high-resolution wall-clock timing to accurately measure and compare execution time for serial and parallel blocks.

- **What scheduling type is used in this program's for loop, and why?**

The default static scheduling is applied here. It evenly divides iterations among threads, which works well when each iteration has nearly equal computational load, as in this case.

- **What is the purpose of `omp_set_num_threads(4)`?**

It explicitly sets the number of OpenMP threads to 4 for the parallel region. This helps control resource usage and test performance scalability.

- **How does the reduction clause avoid race conditions?**

Each thread maintains a private copy of the variable (`parallel_count`), and after the loop, these are combined in a thread-safe manner into a single final result.

- **Could this program benefit from dynamic scheduling? Why or why not?**

Not significantly, because the computation per iteration (checking if a number is prime) is relatively uniform. Dynamic scheduling is more beneficial for irregular workloads.

- **What is the observed performance gain when using OpenMP parallelization in this program?**

The performance gain depends on the hardware and input size. Typically, parallel execution reduces time compared to serial execution, especially for large n, by utilizing multiple cores.

5. Write a MPI Program to demonstration of MPI_Send and MPI_Recv.

```
#include <mpi.h>
#include <stdio.h>
#include <string.h>

int main(int argc, char *argv[] )
{
    int rank, size;
    MPI_Status status;

    MPI_Init(&argc, &argv);           // Initialize the MPI environment
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); // Get the rank of the process
    MPI_Comm_size(MPI_COMM_WORLD, &size); // Get the total number of processes

    if (size < 2) {
        if (rank == 0) {
            printf("This program requires at least two processes.\n");
        }
        MPI_Finalize();
        return 0;
    }

    if (rank == 0) {
        char message[] = "Hello from Process 0 to Process 1";
        printf("Process %d sending message: %s\n", rank, message);
        MPI_Send(message, strlen(message) + 1, MPI_CHAR, 1, 0, MPI_COMM_WORLD);
    } else if (rank == 1) {
        char received_message[100];
        MPI_Recv(received_message, 100, MPI_CHAR, 0, 0, MPI_COMM_WORLD,
        &status);
        printf("Process %d received message: %s\n", rank, received_message);
    }

    MPI_Finalize();
    return 0;
}
```

Step 1: Install MPI : **sudo apt update**

sudo apt install openmpi-bin openmpi-common libopenmpi-dev

Step 2: Check MPI installation:

mpirun –version

Step 3: Save the MPI Program : **mpi_send_recv.c**

Step 4: Compile the MPI Program: **mpicc mpi_send_recv.c -o mpi_send_recv**

Step 5: Run the MPI Program: **mpirun -np 2 ./mpi_send_recv**

Output:

Process 0 sending message: Hello from Process 0 to Process 1

Process 1 received message: Hello from Process 0 to Process 1

Viva Questions:

- **What is the purpose of MPI_Send and MPI_Recv?**

They are used for point-to-point communication to send and receive messages between processes in MPI.

- **What does the tag parameter signify in MPI_Send and MPI_Recv?**

It identifies the message and helps match send and receive operations correctly.

- **What is MPI_COMM_WORLD?**

It is the default communicator that includes all MPI processes in a given program.

- **What is the role of MPI_Status?**

It provides information about a received message, such as the source and tag.

- **What is the output if the program is run with only one process?**

It prints a message that at least two processes are required and exits.

- **Why do we use strlen(message) + 1 in MPI_Send?**

To include the null-terminator \0 so the string can be correctly received and printed.

- **Can MPI_Send and MPI_Recv be used in a non-blocking way?**

Not directly—non-blocking communication uses MPI_Isend and MPI_Irecv.

- **What happens if the receiver buffer size is smaller than the message?**

It can cause a buffer overflow or program crash due to insufficient memory.

- **What ensures synchronization between sender and receiver?**

Blocking behavior of MPI_Send and MPI_Recv ensures synchronization.

- **Can multiple processes use the same tag for different messages?**

Yes, as long as the source and tag combination is correctly matched.

- 6. Write a MPI program to demonstration of deadlock using point to point communication and avoidance of deadlock by altering the call sequence.**

```
#include <mpi.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
int main(int argc, char *argv[]) {
    int rank, size;
    int msg_send = 100, msg_recv;
    MPI_Status status;
    int cause_deadlock = 0;
    // Initialize MPI
    MPI_Init(&argc, &argv);
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); // Get process rank
    MPI_Comm_size(MPI_COMM_WORLD, &size); // Get number of processes
    // Check for mode argument
    if (argc != 2) {
        if (rank == 0)
            printf("Usage: %s --deadlock | --avoid\n", argv[0]);
        MPI_Finalize();
        return 1;
    }
    if (strcmp(argv[1], "--deadlock") == 0) {
        cause_deadlock = 1;
    } else if (strcmp(argv[1], "--avoid") == 0) {
        cause_deadlock = 0;
    } else {
        if (rank == 0)
            printf("Invalid argument. Use --deadlock or --avoid\n");
        MPI_Finalize();
        return 1;
    }
    if (size != 2) {
        if (rank == 0)
            printf("This program requires exactly 2 processes.\n");
        MPI_Finalize();
        return 1;
    }
    // Begin communication
    if (cause_deadlock) {
        // DEADLOCK: both send first, then receive
        if (rank == 0) {
            printf("Process 0 sending to Process 1...\n");
            MPI_Send(&msg_send, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
            printf("Process 0 receiving from Process 1...\n");
            MPI_Recv(&msg_recv, 1, MPI_INT, 1, 0, MPI_COMM_WORLD, &status);
            printf("Process 0 received from Process 1: %d\n", msg_recv);
        } else if (rank == 1) {
            printf("Process 1 sending to Process 0...\n");
        }
    }
}
```

```
MPI_Send(&msg_send, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);
printf("Process 1 receiving from Process 0...\n");
MPI_Recv(&msg_recv, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
printf("Process 1 received from Process 0: %d\n", msg_recv);
}
} else {
// DEADLOCK AVOIDANCE: rank 0 sends first, rank 1 receives first
if (rank == 0) {
printf("Process 0 sending to Process 1...\n");
MPI_Send(&msg_send, 1, MPI_INT, 1, 0, MPI_COMM_WORLD);
printf("Process 0 receiving from Process 1...\n");
MPI_Recv(&msg_recv, 1, MPI_INT, 1, 0, MPI_COMM_WORLD, &status);
printf("Process 0 received from Process 1: %d\n", msg_recv);
} else if (rank == 1) {
printf("Process 1 receiving from Process 0...\n");
MPI_Recv(&msg_recv, 1, MPI_INT, 0, 0, MPI_COMM_WORLD, &status);
printf("Process 1 received from Process 0: %d\n", msg_recv);
printf("Process 1 sending to Process 0...\n");
MPI_Send(&msg_send, 1, MPI_INT, 0, 0, MPI_COMM_WORLD);
}
}
MPI_Finalize();
return 0;
}
```

Step 1: Save the MPI Program : `mpi_deadlock_demo.c`

Step 2: Compile the MPI Program:

`mpicc mpi_deadlock_demo.c -o mpi_deadlock_demo`

Step 3: Run the MPI Program: `mpirun -np 2 ./mpi_deadlock_demo --deadlock`

Step 4: Run the MPI Program: `mpirun -np 2 ./mpi_deadlock_demo --avoid`

Output 1:

Process 0 sending to Process 1...

Process 1 sending to Process 0...

Output 2:

Process 0 sending to Process 1...

Process 1 receiving from Process 0...

Process 1 received from Process 0: 100

Process 1 sending to Process 0...

Process 0 receiving from Process 1...

Process 0 received from Process 1: 100

Viva Questions:

- **What is deadlock in MPI?**

Deadlock occurs when two or more processes wait indefinitely for each other to complete communication.

- **What MPI functions are used in point-to-point communication?**

`MPI_Send()` and `MPI_Recv()`.

- **How can deadlock occur using MPI_Send() and MPI_Recv()?**

If both processes use blocking `MPI_Send()` before `MPI_Recv()`, they can wait indefinitely.

- **How can deadlock be avoided in point-to-point communication?**

By changing the order of send/receive operations or using non-blocking functions like `MPI_Isend()` and `MPI_Irecv()`.

- **What is a blocking communication?**

The function waits until the operation completes before moving forward.

- **What is a tag in MPI communication?**

A message identifier used to match sends and receives.

- **Why is it necessary to use MPI_COMM_WORLD?**

It defines the default communicator including all processes.

- **What is MPI_STATUS_IGNORE used for?**

It ignores the status return of a receive operation.

- **What is the role of MPI_Init and MPI_Finalize?**

They initialize and clean up the MPI environment.

- **How many processes are required to demonstrate deadlock?**

At least two.

7. Write a MPI Program to demonstration of Broadcast operation.

```
#include <stdio.h>
#include <mpi.h>
int main(int argc, char *argv[]) {
    int rank, size;
    int data;

    // Initialize the MPI environment
    MPI_Init(&argc, &argv);

    // Get the rank (ID) of the current process
    MPI_Comm_rank(MPI_COMM_WORLD, &rank);

    // Get the total number of processes
    MPI_Comm_size(MPI_COMM_WORLD, &size);

    if (rank == 0) {
        // Root process initializes the data
        data = 42;
        printf("Process %d (Root) broadcasting data = %d\n", rank, data);
    }

    // Broadcast the data from root (process 0) to all processes
    MPI_Bcast(&data, 1, MPI_INT, 0, MPI_COMM_WORLD);

    // All processes (including root) receive the broadcasted data
    printf("Process %d received data = %d\n", rank, data);

    // Finalize the MPI environment
    MPI_Finalize();
    return 0;
}
```

Step 1: Save the MPI Program : `mpi_broadcast_demo.c`

Step 2: Compile the MPI Program:

```
mpicc mpi_broadcast_demo.c -o mpi_broadcast_demo
```

Step 3: Run the MPI Program: `mpirun -np 4 ./mpi_broadcast_demo`

Output:

```
Process 0 (Root) broadcasting data = 42
Process 0 received data = 42
Process 1 received data = 42
Process 2 received data = 42
Process 3 received data = 42
```

Viva Questions:

- **What is the purpose of MPI_Bcast()?**
It broadcasts data from one process (root) to all other processes.
- **Which process initiates the broadcast?**
The root process.
- **Can the root process receive the data as well?**
Yes, it also participates in the broadcast.
- **What kind of data can be broadcasted?**
Any datatype supported by MPI, such as MPI_INT, MPI_FLOAT, etc.
- **What arguments are required by MPI_Bcast()?**
Buffer, count, datatype, root rank, and communicator.
- **Does MPI_Bcast use point-to-point internally?**
Yes, it is implemented using multiple send/receive operations internally.
- **Is MPI_Bcast blocking?**
Yes, it blocks until the broadcast is complete.
- **How many times is MPI_Bcast called in a program?**
Once by each process for a broadcast operation.
- **Can we have multiple broadcasts in one program?**
Yes.
- **What happens if a non-root process calls MPI_Bcast with a different buffer size?**
It leads to undefined behavior or runtime errors.

8. Write a MPI Program demonstration of MPI_Scatter and MPI_Gather.

```
#include <stdio.h>
#include <mpi.h>
int main(int argc, char *argv[]) {
    int rank, size;
    int send_data[100]; // Array for root to scatter
    int recv_data; // Each process receives one element
    int gathered_data[100]; // Root will gather data here
    MPI_Init(&argc, &argv); // Initialize MPI
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); // Get process rank
    MPI_Comm_size(MPI_COMM_WORLD, &size); // Get total processes
    if (rank == 0) {
        // Initialize data to scatter (only root does this)
        for (int i = 0; i < size; i++) {
            send_data[i] = i * 10; // Example data: 0, 10, 20, ...
        }
        printf("Process 0 (Root): Data prepared for scattering:\n");
        for (int i = 0; i < size; i++) {
            printf("%d ", send_data[i]);
        }
        printf("\n");
    }
    // Scatter: each process receives one element of the send_data array
    MPI_Scatter(send_data, 1, MPI_INT, &recv_data, 1, MPI_INT, 0, MPI_COMM_WORLD);

    // Each process prints what it received
    printf("Process %d received value: %d\n", rank, recv_data);

    // Each process modifies its data (optional)
    recv_data += 1;

    // Gather: all modified recv_data values are sent back to root
    MPI_Gather(&recv_data, 1, MPI_INT, gathered_data, 1, MPI_INT, 0,
    MPI_COMM_WORLD);

    // Root process prints the gathered data
    if (rank == 0) {
        printf("Process 0 (Root): Data gathered after modification:\n");
        for (int i = 0; i < size; i++) {
            printf("%d ", gathered_data[i]);
        }
        printf("\n");
    }
    MPI_Finalize(); // Finalize MPI
    return 0;
}
```

Step 1: Save the MPI Program : `mpi_scatter_gather_demo.c`

Step 2: Compile the MPI Program:

`mpicc mpi_scatter_gather_demo.c -o mpi_scatter_gather_demo`

Step 3: Run the MPI Program: `mpirun -np 4 ./mpi_scatter_gather_demo`

Output:

Process 0 (Root): Data prepared for scattering: 0 10 20 30

Process 0 received value: 0

Process 1 received value: 10

Process 2 received value: 20

Process 3 received value: 30

Process 0 (Root): Data gathered after modification: 1 11 21 31

Viva Questions:

- **What does MPI_Scatter do?**
It divides data from the root and sends portions to all processes.
- **What is MPI_Gather used for?**
It collects data from all processes and assembles it at the root.
- **Do all processes need to call MPI_Scatter/MPI_Gather?**
Yes, all processes in the communicator must call them.
- **Is the data distributed equally in MPI_Scatter?**
Yes, each process receives an equal portion of the total data.
- **What if the number of elements is not divisible by the number of processes?**
It may cause errors or incomplete data distribution.
- **What is the role of the root in MPI_Scatter and MPI_Gather?**
The root sends and receives the complete data buffer, respectively.
- **Are MPI_Scatter and MPI_Gather blocking?**
Yes, both are blocking operations.
- **What is the signature of MPI_Scatter?**
`MPI_Scatter(sendbuf, sendcount, sendtype, recvbuf, recvcount, recvtype, root, comm)`
- **Can these operations be performed without the root process?**
No, the root is mandatory for data distribution and collection.
- **Can we use MPI_Scatter with different datatypes?**
No, all processes must use the same datatype.

9. Write a MPI Program to demonstration of MPI_Reduce and MPI_Allreduce (MPI_MAX, MPI_MIN, MPI_SUM, MPI_PROD).

```
#include <mpi.h>
#include <stdio.h>
int main(int argc, char* argv[]) {
    int rank, size;
    int value, sum, product, max, min;
    int all_sum, all_product, all_max, all_min;
    MPI_Init(&argc, &argv); // Initialize the MPI environment
    MPI_Comm_rank(MPI_COMM_WORLD, &rank); // Get the rank of the process
    MPI_Comm_size(MPI_COMM_WORLD, &size); // Get the total number of processes

    // Each process has its own value (for simplicity, we use rank + 1)
    value = rank + 1;
    printf("Process %d has value: %d\n", rank, value);

    // =====
    // Reduce operations (results gathered at root process - rank 0)
    // =====

    MPI_Reduce(&value, &sum, 1, MPI_INT, MPI_SUM, 0, MPI_COMM_WORLD);
    MPI_Reduce(&value, &product, 1, MPI_INT, MPI_PROD, 0, MPI_COMM_WORLD);
    MPI_Reduce(&value, &max, 1, MPI_INT, MPI_MAX, 0, MPI_COMM_WORLD);
    MPI_Reduce(&value, &min, 1, MPI_INT, MPI_MIN, 0, MPI_COMM_WORLD);
    if (rank == 0) {
        printf("\n--- MPI_Reduce Results (at Root Process 0) ---\n");
        printf("Sum = %d\n", sum);
        printf("Product = %d\n", product);
        printf("Maximum = %d\n", max);
        printf("Minimum = %d\n", min);
    }

    // =====
    // Allreduce operations (results available at all processes)
    // =====

    MPI_Allreduce(&value, &all_sum, 1, MPI_INT, MPI_SUM, MPI_COMM_WORLD);
    MPI_Allreduce(&value, &all_product, 1, MPI_INT, MPI_PROD, MPI_COMM_WORLD);
    MPI_Allreduce(&value, &all_max, 1, MPI_INT, MPI_MAX, MPI_COMM_WORLD);
    MPI_Allreduce(&value, &all_min, 1, MPI_INT, MPI_MIN, MPI_COMM_WORLD);
    printf("\nProcess %d - MPI_Allreduce Results:\n", rank);
    printf("Sum = %d\n", all_sum);
    printf("Product = %d\n", all_product);
    printf("Maximum = %d\n", all_max);
    printf("Minimum = %d\n", all_min);
    MPI_Finalize();
    return 0;
}
```

Step 1: Save the MPI Program : `mpi_reduce_allreduce.c`

Step 2: Compile the MPI Program:

`mpicc mpi_reduce_allreduce.c -o mpi_reduce_allreduce`

Step 3: Run the MPI Program: `mpirun -np 4 ./mpi_reduce_allreduce`

Output:

Process 0 has value: 1

Process 1 has value: 2

Process 2 has value: 3

Process 3 has value: 4

--- MPI_Reduce Results (at Root Process 0) ---

Sum = 10

Product = 24

Maximum = 4

Minimum = 1

Process 0 - MPI_Allreduce Results:

Sum = 10

Product = 24

Maximum = 4

Minimum = 1

Process 1 - MPI_Allreduce Results:

Sum = 10

Product = 24

Maximum = 4

Minimum = 1

Process 2 - MPI_Allreduce Results:

Sum = 10

Product = 24

Maximum = 4

Minimum = 1

Process 3 - MPI_Allreduce Results:

Sum = 10

Product = 24

Maximum = 4

Minimum = 1

Viva Questions:

- **What is the purpose of MPI_Reduce?**
It performs a reduction operation (e.g., sum, max) and returns the result to the root.
- **What does MPI_Allreduce do?**
It performs a reduction and distributes the result to all processes.
- **What are some common operations in MPI_Reduce?**
MPI_SUM, MPI_MAX, MPI_MIN, MPI_PROD.
- **How does MPI_Reduce differ from MPI_Allreduce?**
MPI_Reduce returns result to root; MPI_Allreduce to all.
- **Are these reduction operations commutative?**
Yes, MPI reduction operations are assumed to be commutative and associative.
- **Can we define custom reduction operations?**
Yes, using MPI_Op_create.
- **What is the significance of the root parameter in MPI_Reduce?**
It specifies which process receives the final result.
- **Can MPI_Allreduce be used without a root?**
Yes, it returns results to all processes.
- **Are these operations synchronous?**
Yes, they are blocking and require participation from all processes.
- **What happens if datatypes mismatch in reduction?**
The program leads to undefined behavior or runtime errors.